Analysis of Neural Latent Drift in Large Language Models

Jane Doe, John Smith

Abstract
This paper investigates representation stability during language model fine-tuning.

1. Introduction
Large language models exhibit drift.

References
[1] A. Vaswani et al., Attention is all you need, NeurIPS 2017.